

```
In [1]: import json
import numpy as np
import pandas as pd
import tensorflow as tf
```

```
In [2]: import wget
url='https://storage.googleapis.com/laurencemoroney-blog.appspot.com/sarcasm.json'
filename = wget.download(url)
filename
```

```
100% [.....] 5643545 / 5643545
```

```
Out[2]: 'sarcasm (1).json'
```

```
In [3]: data=pd.read_json(filename)
data.head()
```

```
Out[3]:
```

	article_link	headline	is_sarcastic
0	https://www.huffingtonpost.com/entry/versace-b...	former versace store clerk sues over secret 'b...	0
1	https://www.huffingtonpost.com/entry/roseanne-...	the 'roseanne' revival catches up to our thorn...	0
2	https://local.theonion.com/mom-starting-to-fea...	mom starting to fear son's web series closest ...	1
3	https://politics.theonion.com/boehner-just-wan...	boehner just wants wife to listen, not come up...	1
4	https://www.huffingtonpost.com/entry/jk-rowlin...	j.k. rowling wishes snape happy birthday in th...	0

```
In [4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26709 entries, 0 to 26708
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   article_link    26709 non-null  object
1   headline        26709 non-null  object
2   is_sarcastic    26709 non-null  int64
dtypes: int64(1), object(2)
memory usage: 626.1+ KB
```

```
In [5]: df = data.iloc[:,1:]
df.head()
```

```
Out[5]:
```

	headline	is_sarcastic
0	former versace store clerk sues over secret 'b...	0
1	the 'roseanne' revival catches up to our thorn...	0
2	mom starting to fear son's web series closest ...	1
3	boehner just wants wife to listen, not come up...	1
4	j.k. rowling wishes snape happy birthday in th...	0

```
df.shape
```

In [6]:

Out[6]: (26709, 2)

```
In [7]: sentence = df.headline  
labels = df.is_sarcastic
```

```
In [8]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(sentence, labels, train_size=0.80, random_state=2)
```

```
In [9]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[9]: ((21367,), (5342,), (21367,), (5342,))

```
In [10]: from tensorflow.keras.preprocessing.text import Tokenizer  
tokenizer = Tokenizer(num_words=10000, oov_token="<OOV>")  
tokenizer.fit_on_texts(X_train)
```

```
In [11]: word_index = tokenizer.word_index  
len(word_index)
```

Out[11]: 26527

```
In [12]: training_sequences = tokenizer.texts_to_sequences(X_train)  
training_sequences[0]
```

Out[12]: [42, 103, 87, 69, 158, 94, 167, 170, 1, 23, 342, 4452]

```
In [13]: from tensorflow.keras.preprocessing.sequence import pad_sequences  
  
training_padded = pad_sequences(training_sequences,  
                                maxlen=100,  
                                padding='post',  
                                truncating='post')  
  
training_padded[0]
```

```
Out[13]: array([[ 42, 103, 87, 69, 158, 94, 167, 170, 1, 23, 342,
          4452, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0])
```

```
In [14]: testing_sequences = tokenizer.texts_to_sequences(X_test)

testing_padded = pad_sequences(testing_sequences,
                               maxlen=100,
                               padding='post',
                               truncating='post')
```

```
In [15]: testing_padded[0]
```

```
Out[15]: array([[2496, 30, 4, 12, 27, 1, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           0])
```

```
In [16]: #Layer Normalization
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(10000, 16, input_length=100),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
In [17]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 100, 16)	160000
global_average_pooling1d (GlobalAveragePooling1D)	(None, 16)	0
dense (Dense)	(None, 24)	408
dense_1 (Dense)	(None, 1)	25
=====		
Total params: 160,433		
Trainable params: 160,433		
Non-trainable params: 0		

```
In [18]: history = model.fit(training_padded,
                             y_train,
                             epochs=10,
                             validation_data=(testing_padded, y_test),
                             verbose=2)
```

```

Epoch 1/10
668/668 - 3s - loss: 0.6633 - accuracy: 0.5910 - val_loss: 0.5654 - val_accuracy: 0.8162 - 3s/epoch - 5ms/step
Epoch 2/10
668/668 - 3s - loss: 0.4194 - accuracy: 0.8327 - val_loss: 0.3668 - val_accuracy: 0.8454 - 3s/epoch - 5ms/step
Epoch 3/10
668/668 - 3s - loss: 0.3069 - accuracy: 0.8769 - val_loss: 0.3385 - val_accuracy: 0.8519 - 3s/epoch - 5ms/step
Epoch 4/10
668/668 - 3s - loss: 0.2591 - accuracy: 0.8968 - val_loss: 0.3226 - val_accuracy: 0.8624 - 3s/epoch - 4ms/step
Epoch 5/10
668/668 - 3s - loss: 0.2255 - accuracy: 0.9119 - val_loss: 0.3204 - val_accuracy: 0.8663 - 3s/epoch - 4ms/step
Epoch 6/10
668/668 - 3s - loss: 0.2014 - accuracy: 0.9236 - val_loss: 0.3246 - val_accuracy: 0.8663 - 3s/epoch - 5ms/step
Epoch 7/10
668/668 - 3s - loss: 0.1806 - accuracy: 0.9328 - val_loss: 0.3326 - val_accuracy: 0.8658 - 3s/epoch - 5ms/step
Epoch 8/10
668/668 - 3s - loss: 0.1631 - accuracy: 0.9400 - val_loss: 0.3419 - val_accuracy: 0.8632 - 3s/epoch - 5ms/step
Epoch 9/10
668/668 - 3s - loss: 0.1488 - accuracy: 0.9463 - val_loss: 0.3758 - val_accuracy: 0.8553 - 3s/epoch - 5ms/step
Epoch 10/10
668/668 - 3s - loss: 0.1367 - accuracy: 0.9500 - val_loss: 0.3696 - val_accuracy: 0.8618 - 3s/epoch - 5ms/step

```

```
In [19]: pd.DataFrame(history.history).head()
```

```
Out[19]:
```

	loss	accuracy	val_loss	val_accuracy
0	0.663316	0.591005	0.565383	0.816174
1	0.419370	0.832686	0.366813	0.845376
2	0.306855	0.876913	0.338508	0.851928
3	0.259105	0.896757	0.322640	0.862411
4	0.225541	0.911873	0.320427	0.866342

Plot accuracy vs val_accuracy & loss vs val_loss curve



```
In [20]: import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
```



```
plt.ylabel(string)
plt.legend([string, 'val_'+string])
plt.title(f"{string} vs {'val_'+string}", fontsize = 15)
plt.show()
```

```
plot_graphs(history, "accuracy")
plot_graphs(history, "loss")
```



```
In [21]: model.evaluate(testing_padded,y_test)
```

```
167/167 [=====] - 0s 1ms/step - loss: 0.3696 - accuracy: 0.8618  
Out[21]: [0.3695674240589142, 0.8618494868278503]
```

Predictions

```
In [22]: pred = model.predict(testing_padded)  
pred[:5]
```

```
Out[22]: array([[0.20597145],  
                [0.8800808 ],  
                [0.35702163],  
                [0.1383534 ],  
                [0.99725986]], dtype=float32)
```

```
In [23]: pred_labels= [1 * (i[0]>=0.5) for i in pred]  
pred_labels[:5]
```


```
Out[23]: [0, 1, 0, 0, 1]
```

```
In [24]: y_test.head(5)
```

```
Out[24]: 10465    0  
        970     1  
        17936    1  
        9001     0  
        12922    1  
        Name: is_sarcastic, dtype: int64
```

```
In [25]: #Confusion_Matrix = pd.crosstab(y_test, pred_labels)  
Confusion_Matrix=tf.math.confusion_matrix(labels = y_test, predictions= pred_labels)
```

```
In [26]: import seaborn as sns  
sns.heatmap(Confusion_Matrix, annot = True, fmt='d')  
plt.xlabel('Predicted', fontsize=13)  
plt.ylabel('Truth', fontsize=13)  
plt.title('Confusion Matrix', fontsize=17)  
plt.show()
```

```
In [27]: testing_padded[0]
```

[illegible]

```
In [28]: # Reverse the word_index
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
reverse_word_index[1084]
```

```
Out[28]: 'announce'
```

```
In [29]: # Decode the sentence
def decode_sentence(text):
    word=list()
    for i in text:
        if i != 0:
            word.append(reverse_word_index[i])
```

```
sentence = ' '.join(word)
return 'sentence :', sentence
```

```
In [30]: decode_sentence(testing_padded[0])
```

```
Out[30]: ('sentence :', 'glasses are the new it <00V>')
```

```
In [31]: y_test.head(1)
```

```
Out[31]: 10465    0
         Name: is_sarcastic, dtype: int64
```

```
In [32]: e = model.layers[0]
         e
```

```
Out[32]: <keras.layers.embeddings.Embedding at 0x2b5be807820>
```

```
In [33]: weights = e.get_weights()[0]
         weights[0]
```

```
Out[33]: array([ 0.1600546 ,  0.01401678,  0.00055776,  0.04858763,  0.0108338 ,
                -0.12296436, -0.0004041 ,  0.3463165 ,  0.03926046, -0.01363823,
                -0.04651956, -0.03632296, -0.00501014, -0.05600563,  0.07940689,
                -0.0232128 ], dtype=float32)
```

```
In [34]: vocab_size, embedding_dim = weights.shape
         vocab_size
```

```
Out[34]: 10000
```

Save tokenize set & Predicted value

```
In [35]: import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
for word_num in range(1, vocab_size):
    word = reverse_word_index[word_num]
    embeddings = weights[word_num]
    out_m.write(word + "\n")
```

```
out_v.write('\t'.join([str(x) for x in embeddings]) + "\n")
out_v.close()
out_m.close()
```

```
In [36]: sentence = ["No matter who wins, This election will seriously boost alcohol sales."]
sequences = tokenizer.texts_to_sequences(sentence)
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')
np.argmax(model.predict(padded))
```

Out[36]: 0

```
In [37]: sentence = ["I work 40 hours a week for me to be this poor!"]
sequences = tokenizer.texts_to_sequences(sentence)
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')
np.argmax(model.predict(padded))
```

Out[37]: 0

```
In [38]: sentence = ["Well, what a surprise."]
sequences = tokenizer.texts_to_sequences(sentence)
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')
np.argmax(model.predict(padded))
```

Out[38]: 0

```
In [39]: sentence = ['I will consider myself lucky to die alone']
sequences = tokenizer.texts_to_sequences(sentence)
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')
np.argmax(model.predict(padded))
```

Out[39]: 0

```
In [40]: sentence = ["Friday is my second favorite f-word"]
sequences = tokenizer.texts_to_sequences(sentence)
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')
np.argmax(model.predict(padded))
```

Out[40]: 0

```
In [41]: sentence = ['Hope followers have increased!']
sequences = tokenizer.texts_to_sequences(sentence)
```

```
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')  
np.argmax(model.predict(padded))
```

Out[41]: 0

```
In [42]: sentence = ['So let it be with Caesar.  
The noble Brutus / Hath told you Caesar was ambitious.  
If it were so, it was a grievous fault,  
And grievously hath Caesar answered it.']  
sequences = tokenizer.texts_to_sequences(sentence)  
padded = pad_sequences(sequences, maxlen=100, padding='post', truncating='post')  
np.argmax(model.predict(padded))
```

Out[42]: 0